Understanding the Rational Approximation of the Exponential Integrator (REXI)

Martin Schreiber <M.Schreiber@exeter.ac.uk>
Pedro S. Peixoto <pedrosp@ime.usp.br>
et al.

September 16, 2015

This document serves as the basis for implementing the Rational approximation of the EXponential Integrator (REXI). Here, we purely focus on the linear part of the shallow-water equations (SWE) and show the different steps to approximate solving this linear part with an exponential integrator. This paper mainly summarises previous work on REXI.

1 Problem formulation

We use linearised shallow water equations (SWE) with respect to a rest state with mean water depth of η_0 and defined for perturbations of height η (see [1]). The linear operator (L) may be written as

$$L(U) := \begin{pmatrix} 0 & -\eta_0\partial_x & -\eta_0\partial_y \\ -g\partial_x & 0 & f \\ -g\partial_y & -f & 0 \end{pmatrix} U$$

where $U := (\eta, u, v)^T$. Here, we neglect all non-linear terms and consider f constant (f-plane approximation). The time evolution of the PDE, with the subscript t denoting the derivative in time, is given by

$$U_t = L(U).$$

It is further worth noting, that this system describes an oscillatory system (2D wave equation), hence the operator L is hyperbolic and has imaginary eigenvalues. Linear initial value differential problems are well known to be solvable with exponential integrators for arbitrary time step sizes via

$$U(t) = e^{Lt}U(0).$$

see e.g. [3]. However, this is typically quite expensive to compute and analytic solutions only exist for some simplified system of equations, see e.g. [1]

1

for f-plane shallow-water equations. These exponential integrators can be approximated with rational functions and this paper is on giving insight into this approximation.

2   1D rational approximation

Terry et al. [2] developed a rational approximation of the exponential integrator. First, we like to get more insight into it with a one-dimensional formulation before applying REXI to a rational approximation of a linear operator. Our main target is to find an approximation of an operator with a *complex exponential shape*, in our case e^{ix}, which (in one-dimension) is given as a function $f(x)$. We will end up in an approximation given by the following rational approximation:

$$e^{ix} \approx \sum_{n=-N}^{N} Re\left(\frac{\beta_n}{ix + \alpha_n}\right)$$

with complex coefficients α_n and β_n. We point out that the coefficients α_n will always have non zero real part, so no singularity occurs with the rational function.

2.1   Step A: Approximation of solution space

First, we assume that we can use Gaussian curves as basis functions for our approximation. So first we find an approximation of one of our underlying Gaussian basis function

$$\psi_h(x) := (4\pi)^{-\frac{1}{2}} e^{-x^2/(4h^2)}$$

In this formulation, h can be interpreted as the horizontal "stretching" of the basis function. Note the similarities to the Gaussian distribution, but by dropping certain parts of the vertical scaling as it is required for probability distributions. We can now approximate our function $f(x)$ with a superposition of basis functions $\psi_h(x)$ by

$$f(x) \approx \sum_{m=-M}^{M} b_m \psi_h(x + mh)$$

with M controlling the interval of approximation (~size of "domain of interest") and h will be related to the accuracy of integration (~resolution in "domain of interest"). We choose h small enough so that the support of the Fourier transform of f is mainly localised within $[-1/(2h), 1/(2h)]$, i.e. almost zero outside this interval. M is chosen such that the approximation will be adequate in the interval $|x| < Mh$.

To compute the coefficients b_m, we rewrite the previous equation in Fourier space with

$$\frac{\hat{f}(\xi)}{\hat{\psi}_h(\xi)} = \sum_{m=\infty}^{\infty} b_m e^{2\pi i m h \xi},$$

where the $\hat{\ }$ symbols indicate the Fourier transforms of the respective functions. The b_m are now the Fourier coefficients of the series for the function $\frac{\hat{f}(\xi)}{\hat{\psi}_h(\xi)}$ and can be calculated as (see [2], page 11),

$$b_m = h \int_{-\frac{1}{2h}}^{\frac{1}{2h}} e^{-2\pi i m h \xi} \frac{\hat{f}(\xi)}{\hat{\psi}_h(\xi)} d\xi,$$

for m integer and $1/h$ defines the periodicity of the trigonometric basis function.

Since we are interested in approximating $f(x) = e^{ix}$, we can simplify the equation by using the response in frequency space $\hat{f}(\xi) = \delta(\xi - \frac{1}{2\pi})$, where here δ is the Dirac distribution, and

$$b_m = h\, e^{-imh} \hat{\psi}_h \left(\frac{1}{2\pi} \right)^{-1}.$$

The Fourier transform of the Gaussian function is well known and given by

$$
\begin{aligned}
\hat{\psi}_h(\xi) &= \int_{-\infty}^{\infty} \frac{1}{\sqrt{4\pi}} e^{-\left(\frac{x}{2h}\right)^2} e^{-2\pi i x \xi} dx \\
&= \frac{1}{\sqrt{4\pi}} \int_{-\infty}^{\infty} e^{-\left(\left(\frac{x}{2h}\right)^2 + 2\pi i x \xi + (2h\pi i \xi)^2 - (2h\pi i \xi)^2\right)} dx \\
&= \frac{1}{\sqrt{4\pi}} e^{-(2h\pi \xi)^2} \int_{-\infty}^{\infty} e^{-\left(\frac{x}{2h} + 2h\pi i \xi\right)^2} dx \\
&= \frac{1}{\sqrt{4\pi}} e^{-(2h\pi \xi)^2} \int_{-\infty}^{\infty} e^{-\left(\frac{x}{2h}\right)^2} dx \\
&= h e^{-(2h\pi \xi)^2}
\end{aligned}
$$

where we used that $\int_{-\infty}^{\infty} e^{-\left(\frac{x}{2h}\right)^2} dx = h\sqrt{4\pi}$ and completed squares in the exponential term. For the case $\xi = \frac{1}{2\pi}$, we get

$$\hat{\psi}_h \left(\frac{1}{2\pi} \right) = h\, e^{-h^2}.$$

Finally, one can obtain the equation

$$b_m = h\,e^{-imh}\frac{1}{h\,e^{-h^2}} = e^{-imh}e^{h^2}$$

to compute the coefficients b_m for $f(x) = e^{ix}$.

2.2 Step B: Approximation of basis function

The second step is the approximation of the basis function $\psi_h(x)$ itself with a rational approximation, see [4]. Our basis function is given by

$$\psi_h(x) := (4\pi)^{-\frac{1}{2}}e^{-x^2/(4h^2)}$$

and a close-to-optimal approximation of $\psi_1(x)$ with a sum of rational functions is given by

$$\psi_1(x) \approx Re\left(\sum_{l=-L}^{L}\frac{a_l}{ix + (\mu + i\,l)}\right)$$

with the μ and a_l given in [4], Table 1. We can generalise this approximation to arbitrary chosen h via

$$\psi_h(x) \approx Re\left(\sum_{l=-L}^{L}\frac{a_l}{i\frac{x}{h} + (\mu + i\,l)}\right).$$

The theory of how these coefficients are calculated are in [4] and will not be described here. Therefore, we assume that the coefficients a_l are given.

2.3 Step C: Approximation of the approximation

We then combine the approximation (B) into the approximation (A), yielding the approximation $\tilde{f}$ for f given by

$$\tilde{f}(x) = \sum_{m=-M}^{M} b_m\psi_h(x + mh) = \sum_{m=-M}^{M} b_m Re\left(\sum_{l=-L}^{L}\frac{a_l}{i\frac{x+mh}{h} + (\mu + i\,l)}\right)$$

$$= \sum_{m=-M}^{M} b_m \sum_{l=-L}^{L} Re\left(\frac{ha_l}{ix + h(\mu + i(m+l))}\right).$$

We further like to simplify this equation and we observe, that for $n := m+l$, the denominator is equal. We can hence express parts of the denominator in terms of $n := m + l$ by

$$\alpha_n := s(\mu + in). \tag{1}$$

Now, we merge the b_m and a_l coefficients and first have a look at the b_m which is complex valued. We observe the following property: Assuming that we want

to compute the real value of $f(x)$, only the real value of b_m has to be merged with the sum, since the imaginary component would be dropped afterwards. This allows us to move the $Re(b_m)$ values inside the $\sum_L$:

$$Re(f(x)) := Re\left(\sum_{m=-M}^{M} \sum_{l=-L}^{L} \frac{Re(b_m)\, ha_l}{ix + h(\mu + i(m+l))}\right).$$

Now we can collect all nominators with equivalent denominator (if $n = m + l$ and by using δ as the Kronecker delta), yielding

$$\beta_n^{Re} := h \sum_{m=-M}^{M} \sum_{l=-L}^{L} Re(b_m)ha_l\delta(n,\, m+l) \tag{2}$$

for real values $f(x)$ and

$$\beta_n^{Im} := h \sum_{m=-M}^{M} \sum_{l=-L}^{L} Im(b_m)a_l\delta(n,\, m+l) \tag{3}$$

for complex values of $f(x)$. Note, that the complexity of this operation is $O(NL)$, which is negligible for small N and L. This can be optimized by using $L_1 = max(-L, n - M)$ and $L_2 = min(L, n + M)$ (see [2]) and we can compute β_n^{Re} with

$$\beta_n^{Re} := h \sum_{k=L_1}^{L_2} Re(b_{n-k})a_k, \tag{4}$$

and β_n^{Im} correspondingly.

Finally, we get the REXI approximation

$$e^{ix} \approx \sum_{n=-N}^{N} Re\left(\frac{\beta_n^{Re}}{ix + \alpha_n}\right) + i\, Re\left(\frac{\beta_n^{Im}}{ix + \alpha_n}\right).$$

2.4 Example coefficents

For a better understanding and discussion of the poles, we provide some explanatory coefficients for α_n and β_n^{Re} computed with $M := 2$ and $h := 0.2$:

n	α_n	β_n
-13	(-0.863064302175, -2.6)	(-2.0794560075645e-08,5.312368394177e-09)
-12	(-0.863064302175, -2.4)	(-1.8562925598646e-08,-1.6892470811809e-07)
-11	(-0.863064302175, -2.2)	(6.8570271350932e-07,-4.4377515257134e-08)
-10	(-0.863064302175, -2)	(1.9470768200785e-07,2.1186231739561e-06)
-9	(-0.863064302175, -1.8)	(3.037169144916e-06,-3.8007524015554e-06)
-8	(-0.863064302175, -1.6)	(-0.00020292956274934,-9.4793805592883e-05)
-7	(-0.863064302175, -1.4)	(0.00051562027155282,0.0033198141762956)
-6	(-0.863064302175, -1.2)	(0.023802856324805,-0.020097812439831)
-5	(-0.863064302175, -1)	(-0.16210306892042,-0.057527918763957)
-4	(-0.863064302175, -0.8)	(0.083936569694558,0.55379453117192)
-3	(-0.863064302175, -0.6)	(0.87683903065806,-0.58136186212318)
-2	(-0.863064302175, -0.4)	(-0.87618099667542,-0.6444132979014)
-1	(-0.863064302175, -0.2)	(-0.2112750856805,0.51693268636776)
0	**(-0.863064302175, 0)**	**(0.21113064943379,1.1012434042446e-07)**
1	(-0.863064302175, 0.2)	(-0.2112752777559,-0.51693263772868)
2	(-0.863064302175, 0.4)	(-0.87618105783081,0.6444131761443)
3	(-0.863064302175, 0.6)	(0.87683907406497,0.58136183517238)
4	(-0.863064302175, 0.8)	(0.083936534477108,-0.55379454106338)
5	(-0.863064302175, 1)	(-0.16210304313401,0.057527824955638)
6	(-0.863064302175, 1.2)	(0.023802980792584,0.020097827969804)
7	(-0.863064302175, 1.4)	(0.00051562077173168,-0.0033196934926057)
8	(-0.863064302175, 1.6)	(-0.0002030221163996,9.4802957184526e-05)
9	(-0.863064302175, 1.8)	(3.0281700967037e-06,3.7434363774526e-06)
10	(-0.863064302175, 2)	(2.2311216999616e-07,-2.1234907990132e-06)
11	(-0.863064302175, 2.2)	(6.871098037128e-07,5.5123982746463e-08)
12	(-0.863064302175, 2.4)	(-2.1322288893395e-08,1.6899352278552e-07)
13	(-0.863064302175, 2.6)	(-2.0738399377275e-08,-5.6642992624128e-09)

3 REXI on linear operators

In this section, we investigate the linear operator L with the rational approximation.

3.1 Reducing number of computations for L

Note the property (see Sec. 3.3 in [2]) for the α_n and β_n: There is an anti-symmetry around the central pole with

$$\alpha_{-n} = \bar{\alpha}_{+n} \tag{5}$$

and

$$\beta_{-n} = \bar{\beta}_{+n} \tag{6}$$

In particular, with $Im(\alpha_0) = Im(\beta_0) = 0$, there is a zero imaginary number for the central pole.

Furthermore, it holds that

$$\overline{(L + \alpha)^{-1} U(0)} = (L + \overline{\alpha})^{-1} U(0)$$

with the overbar denoting the complex conjugate. We can then reformulate the approximation

$$
\begin{aligned}
e^{\tau L} &\approx \sum_{n=-N}^{N} Re\left(\frac{\beta_n^{Re}}{\tau L + \alpha_n}\right) \\
&= \sum_{n=-N}^{-1} Re\left(\frac{\beta_n^{Re}}{\tau L + \alpha_n}\right) + Re\left(\frac{\beta_0^{Re}}{\tau L + \alpha_0}\right) + \sum_{n=-N}^{-1} Re\left(\frac{\beta_{N+n+1}^{Re}}{\tau L + \alpha_{N+n+1}}\right)
\end{aligned}
$$

and using the properties (5) and (6), we can write this as

$$
\begin{aligned}
&\sum_{n=-N}^{-1} Re\left(\frac{\beta_n^{Re}}{\tau L + \alpha_n}\right) + Re\left(\frac{\beta_0^{Re}}{\tau L + \alpha_0}\right) + \sum_{n=-N}^{-1} Re\left(\frac{\overline{\beta_N^{Re}}}{\tau L + \overline{\alpha_N}}\right) \\
&= \sum_{n=-N}^{-1}\left(Re\left(\frac{\beta_n^{Re}}{\tau L + \alpha_n} + \overline{\left(\frac{\beta_n^{Re}}{\tau L + \alpha_n}\right)}\right)\right) + Re\left(\frac{\beta_0^{Re}}{\tau L + \alpha_0}\right).
\end{aligned}
$$

(TODO: Is this reformulation in the last line really allowed?!? We are computing the conjugate on τL)
Since the imaginary parts cancel out for $Re(a + \overline{a}) = a + \overline{a}$ and with α_0 and β_0 being only real-valued, we can simplify the equation to

$$e^{\tau L} \approx \sum_{n=-N}^{N} Re\left(\frac{\beta_n^{Re}}{\tau L + \alpha_n}\right) = \sum_{n=-N}^{-1}\left(2\frac{\beta_n^{Re}}{\tau L + \alpha_n}\right) + \frac{\beta_0^{Re}}{\tau L + \alpha_0}.$$

This allows us to reduce the computational amount almost by a factor of two for solving $(L + \alpha)^{-1}$ giving the real valued solution

$$e^{\tau L} \approx \sum_{n=0}^{N} \frac{\gamma_n^{Re}}{\tau L + \alpha_n}$$

with

$$\gamma_n := \begin{cases} \beta_0 & for\, n = 0 \\ 2\beta_n & else \end{cases}$$

3.2   Matrix exponential

We would like to apply REXI to a formulation such as

$$U(t) = e^{\tau L} U(0).$$

To see the relationship between the approximation of e^{ix} with $e^{\tau L}$ we assume that L is skew hermitian and therefore has only purely imaginary eigenvalues, and maybe decomposed as $\Sigma \Lambda \Sigma^H$, yielding

$$e^{\tau L} = \sum_{k=0}^{\infty} \frac{(\tau L)^k}{k!} = \sum_{k=0}^{\infty} \frac{\tau \Sigma \Lambda^k \Sigma^H}{k!} = \Sigma \left(\sum_{k=0}^{\infty} \frac{(\tau \Lambda)^k}{k!} \right) \Sigma^H = \Sigma e^{\tau \Lambda} \Sigma^H, \qquad (7)$$

where we used the orthonormality of Σ to cancel it out from the summation, and

$$e^{\tau \Lambda} = \begin{pmatrix} \ddots & & \\ & e^{i\lambda_j \tau} & \\ & & \ddots \end{pmatrix}$$

where we have explicitly detached the imaginary unit from the eigenvalues, therefore λ_n are assumed real. Since $e^{\tau \Lambda}$ is diagonal, it can be eigenvalue-wise approximated in the same way as in e^{ix} with REXI.

Although L has imaginary eigenvalues, we wish to evaluate the $e^{\tau L} U(0)$, which is real valued. Therefore, we will use the real approximation of e^{ix}

$$exp(\tau L) \approx Re \left(\sum_{n=-N}^{N} \beta_n (\tau L + \alpha_n)^{-1} \right), \qquad (8)$$

where β_n is given by equation (2) and α_n by equation (1).

Given the linear operator L to be applied on $U(0)$, we continue to show the relation of the linear operator to the REXI approximation of the real term

$$e^{\tau L} \approx \sum_{n=-N}^{N} Re \left(\frac{\beta_n^{Re}}{\tau L + \alpha_n} \right)$$

as we use it later on. Next, we rewrite L with our EV decomposition as

$$
\begin{aligned}
e^{\tau L} &\approx \sum_{n=-N}^{N} Re \left(\beta_n^{Re} (\tau \Sigma \Lambda \Sigma^H + \alpha_n I)^{-1} \right) \\
&= \sum_{n=-N}^{N} Re \left(\beta_n^{Re} (\Sigma (\tau \Lambda \Sigma^H + \alpha_n \Sigma^H))^{-1} \right) \\
&= \sum_{n=-N}^{N} Re \left(\beta_n^{Re} (\Sigma (\tau \Lambda + \alpha_n \Sigma^H \Sigma) \Sigma^H)^{-1} \right) \\
&= \sum_{n=-N}^{N} Re \left(\beta_n^{Re} \Sigma^H (\tau \Lambda + \alpha_n)^{-1} \Sigma \right) \\
&= \Sigma^H \left[\sum_{n=-N}^{N} Re \left(\beta_n^{Re} (\tau \Lambda + \alpha_n)^{-1} \right) \right] \Sigma \\
&= \Sigma^H e^{\tau \Lambda} \Sigma \qquad (9)
\end{aligned}
$$

@ALL: equation (7) and (9) should match, but are not.

Note, that we use the same REXI approximation for different λ_j. Hence, the approximation has to be sufficiently accurate over the entire range of all Eigenvaluesλ_j which is what we use the approximation for.

3.3 Choosing h and M

Some important points about the choice of M and h have to be made at this point. We know that e^{ix} is accurately approximated with REXI for the interval $|x| < hM$, where h is chosen small enough to obtain a good approximation in step (A), and M will define the interval size and number of approximation points. In the matrix case, M has to be chosen so that $hM > t\bar{\lambda}$, where $\bar{\lambda} = \max_n |\lambda_n|$, in order to capture all wavelengths of L. In other other words, hM need to be set to capture the fastest wave. Note that if this is used as a time stepping method, with time step $t = \tau$, then, the larger the timestep, the larger M will be. Exact evaluations of the choices for h and M may be done based on equation (3.6) of [2].

TODO: We have to investigate this by far in more depth. Can we possibly use a standard time step restriction of the linear equations?

3.4 Handling τ in REXI

We reformulate the REXI approximation scheme given by

$$(\tau L + \alpha)^{-1} U(\tau) = U(0)$$

and by factoring τ out, yielding

$$(L + \frac{\alpha}{\tau})^{-1} U(\tau)\tau^{-1} = U(0)$$

So instead of solving for $U(\tau)$, we are solving for $U^\tau(\tau) := U(\tau)\tau^{-1}$ as well as $\alpha^\tau := \frac{\alpha}{\tau}$.

To summarize, we have to solve the system of equations given by

$$(L + \alpha^\tau)^{-1} U^\tau(\tau) = U(0) \tag{10}$$

with $U(0)$ the initial conditions. For sake of simplicity, we stick to the formulation without the τ notation.

3.5 Computing inverse of $(L + \alpha)^{-1}$

For computing the inverse, arbitrary solvers can be used. However we like to note, that α is a complex number. Hence, requiring solvers with support for solving in complex space.

We wish to solve the differential problem for each time step

$$(-L + \alpha)U = U(0)$$

so that $U = (-L + \alpha)^{-1} U_0$. Note, the change in sign before L for convenience which has to be accounted for afterwards. We will do this converting the problem into and elliptic equation.

First, lets expand the equations with the definition of L,

$$
\begin{align}
fv + g\eta_x + \alpha u &= u_0, \tag{11} \\
fu + g\eta_y + \alpha v &= v_0, \tag{12} \\
\bar{\eta}(u_x + v_y) + \alpha \eta &= \eta_0. \tag{13}
\end{align}
$$

Let f be constant (f-plane approximation), $\delta := u_x + v_y$ be the wind divergence, $\zeta := v_x - u_y$ be the wind (relative) vorticity and $\Delta \eta := \eta_{xx} + \eta_{yy}$ the Laplacian of the fluid depth. We will re-write the problem in a divergence-vorticity formulation by taking 2 steps. First, sum the ∂_x of equation (11) and the ∂_y of equation (12), yielding

$$
-f\zeta + g\Delta\eta + \alpha\delta = \delta_0. \tag{14}
$$

Then subtract the ∂_y of equation (11) from the ∂_x of equation (12), yielding

$$
f\delta + \alpha\zeta = \zeta_0. \tag{15}
$$

Using equation (14) in equation (15) gives us

$$
\delta = -\frac{1}{f^2 + \alpha^2} \left(\alpha g \Delta\eta - \alpha\delta_0 - f\zeta_0 \right).
$$

Finally, substituting δ in equation (13), that reads $\bar{\eta}\delta + \alpha\eta = \eta_0$, results in

$$
-\frac{\bar{\eta}}{f^2 + \alpha^2} \left(\alpha g \Delta\eta - \alpha\delta_0 - f\zeta_0 \right) + \alpha\eta = \eta_0,
$$

which may be simplified into the elliptic equation by multiplying by $-\frac{f^2 + \alpha^2}{\bar{\eta}\alpha g}$

$$
\Delta\eta - \kappa^2\eta = r_0 \tag{16}
$$

where

$$
\kappa^2 = \frac{f^2 + \alpha^2}{\bar{\eta}g}
$$

and

$$
r_0 = -\frac{\kappa^2}{\alpha}\eta_0 + \frac{1}{g}\delta_0 + \frac{f}{\alpha g}\zeta_0.
$$

Multiplying the equation (16)by $-g\bar{\eta}$ gives

$$
((\alpha^2 + f^2) - g\bar{\eta}\Delta)\eta = \frac{f^2 + \alpha^2}{\alpha}\eta_0 - \bar{\eta}\delta_0 - \frac{f\bar{\eta}}{\alpha}\zeta_0
$$

We can computeη with any elliptic solver. Special attention has to be given to the LHS of the form $(\gamma - \Delta)\eta$ for a spectral solver and this is the reason for

this brief excursion. We use the annotation~to annotate an operator or quantity to be given in spectral space. Then, we factor $\tilde{\eta}$ in and write form of the LHS in spectral space with the identity matrix I as

$$(\gamma I \tilde{\eta} - \tilde{\Delta}\tilde{\eta}) = (\gamma I - \tilde{\Delta})\tilde{\eta}$$

Here, the minus operator in spectral space has to be applied on all elements of $\tilde{\Delta}$ and not only to the 0^{th} mode as it is typically the case for adding a constant in spectral space.

We need to retrieve the velocities by solving the 2×2 system formed by equations (11) and (12), which reads

$$A_\alpha U = U_0 - g\nabla\eta$$

with

$$A_\alpha = \begin{pmatrix} \alpha & -f \\ f & \alpha \end{pmatrix}.$$

The solution is

$$U = A_\alpha^{-1}(U_0 - g\nabla\eta)$$

where

$$A_\alpha^{-1} = \frac{1}{f^2 + \alpha^2} \begin{pmatrix} \alpha & f \\ -f & \alpha \end{pmatrix}$$

Finally, since we computed $(-L + \alpha)^{-1}$, we also have to invert the sign of the computed solution U. Alternatively, we can change the signs of α_n and β_n.

3.6 Interpretation of τ

We like to close this section with a brief discussion of τ by having a look on the REXI reformulation

$$(L - \frac{\alpha}{\tau})^{-1} U(\tau)\tau^{-1} = U(0)$$

We see, that for an increasing τ, hence an integration in time over a larger time period, the poles given by α are getting closer. This can possibly lead to a loss in accuracy for the data sampled by the outer poles α_{-N} and α_N. Therefore, the number N of poles is expected to scale linearly with the size of the coarse time step,

$$|N| \propto \tau.$$

Indeed, we saw in section 3.2 that for larger τ, M needs to be larger.

3.7 Complex solver

We are interested in solving a system of the form,

$$(-L + \alpha)U = U_0,$$

where α is a complex number, therefore we must allow complex solutions for U. The system can be transformed to have only real arithmetic in the following way.

First we decompose U into its real and imaginary parts, $U = U^r + iU^i$, and allow U_0 to be decomposed in the same way. Although α can be a general complex number, it real part is always constant in REXI, given by μ. We will therefore write $\alpha = \mu + i\alpha^i$, and we can absorb μ into L writing $D = -L + \mu I$, where I is the identity matrix. Now

$$(D + i\alpha^i I)(U^r + iU^i) = U_0^r + iU_0^i,$$
$$DU^r - \alpha^i U^i + i(\alpha^i U^r + DU^i) = U_0^r + iU_0^i,$$

therefore

$$DU^r - \alpha^i U^i \quad = \quad U_0^r \tag{17}$$
$$\alpha^i U^r + DU^i \quad = \quad U_0^i, \tag{18}$$

which in matrix notation gives

$$\begin{pmatrix} D & -\alpha^i I \\ \alpha^i I & D \end{pmatrix} \begin{pmatrix} U^r \\ U^i \end{pmatrix} = \begin{pmatrix} U_0^r \\ U_0^i \end{pmatrix},$$

or,

$$\begin{pmatrix} -L + \mu I & -\alpha^i I \\ \alpha^i I & -L + \mu I \end{pmatrix} \begin{pmatrix} U^r \\ U^i \end{pmatrix} = \begin{pmatrix} U_0^r \\ U_0^i \end{pmatrix}.$$

A similar approach may be taken with the elliptic equation

$$\Delta \eta + \theta \eta = r_0, \tag{19}$$

where η, θ and r_0 may be complex. The resulting system is given by

$$\Delta \eta^r + \theta^r \eta^r - \theta^i \eta^i = r_0^r, \tag{20}$$
$$\Delta \eta^i + \theta^i \eta^r + \theta^r \eta^i = r_0^i, \tag{21}$$

which can also be written in matrix notation and solved with arbitrary elliptic equation solvers.

4 Filtering

The method described in the previous section is well defined for skew hermitian L. If L is not skew hermitian, the real eigenvalues might cause the REXI to have absolute values larger than 1, which can lead to instabilities if used as time stepping method.

To ensure that the REXI is bounded by unit, a filtering process is proposed in [2]. REXI is prone to exceed unit in the neighborhood of $|t\lambda| \approx hM$, therefore in the highest frequencies. The idea is to construct a rational function $S(ix)$ that is approximately 1 in a smaller interval $|t\lambda| < hM_0$, with $M_0 < M$, and decays very fast to zero outside this interval. Then we multiply this filters function to the original REXI, which will lead to a unit bounded REXI.

Further details of how $S(ix)$ is computed will be added later.

5   Bringing everything together

Using the spectral methods (e.g. in SWEET), we can directly solve the Helmholtz problem for the height in Eq. (**??**) and then solver for the velocity in Eqs. (**??,??**). Note that the Helmholtz problem is in complex space, as α_n are complex. This is straightforward with spectral methods. For finite difference/element methods, the problem can be split into its real and imaginary parts.

Then, the problem is reduced to computing the REXI as given in Eq. (8).

6   Notes on HPC

- The terms in REXI to solve are all independent. Hence, for latency avoiding, the communication can be interleaved with computations.

- The iterative solvers are memory bound. Instead of computing $c := a * b$ for the stencil operations, we could compute $\vec{c} := a\vec{b}$ with a one coefficient in the stencil. This allows vectorization over c and b on accelerator cards with strided memory access.

- It is unknown which method is more efficient to solve the system of equations:

    - iterative solvers have low memory access,
    - inverting the system and storing it as a sparse matrix allows fast direct solving but can yield more memory access operations.

- Splitting the solver into real and complex number would store them consecutively in memory. This has a potential to avoid non-strided memory access and using the same SIMD operations (Just a rough idea, TODO: check if this is really the case).

7   Acknowledgements

Thanks to Terry for the feedback & discussions!

References

[1] Formulations of the shallow-water equations, M. Schreiber, P. Peixoto et al.

[2] High-order time-parallel approximation of evolution operators, T. Haut et al.

[3] Nineteen Dubious Ways to Compute the Exponential of a Matrix, Twenty-Five Years Later, Cleve Moler and Charles Van Loan, SIAM review

[4] Near optimal rational approximations of large data sets, Damle, A., Beylkin, G., Haut, T. S. & Monzon